

Q1. The decision tree classifier algorithm is a supervised learning algorithm used for classification tasks. It works by recursively partitioning the feature space into regions, based on the values of input features, to create a tree-like structure. At each internal node of the tree, a decision is made based on the value of a specific feature, leading to a split in the data. The goal is to create a tree that accurately predicts the target variable by learning decision rules from the training data.

Q2. The mathematical intuition behind decision tree classification involves recursively partitioning the feature space to minimize impurity or maximize information gain. At each step, the algorithm selects the feature and the split point that best separates the data into purest possible subsets. The impurity of a node is typically measured using metrics like Gini impurity or entropy. The algorithm continues this process until either all leaves are pure or a stopping criterion is met.

Q3. In a binary classification problem, a decision tree classifier can be used to partition the feature space into two regions corresponding to the two classes. At each node of the tree, a decision is made based on the value of a feature, and the data is split into two subsets. The process continues recursively until each leaf node contains instances of only one class, or until a stopping criterion is met.

Q4. Geometrically, decision tree classification can be visualized as dividing the feature space into rectangular regions, where each region corresponds to a different class label. The decision boundaries are axis-parallel, meaning they are perpendicular to the feature axes. To make predictions for a new instance, the algorithm traverses the tree from the root node to a leaf node, following the decision rules based on the feature values.

Q5. The confusion matrix is a table used to evaluate the performance of a classification model. It compares the actual class labels of the data with the predicted class labels produced by the model. The confusion matrix consists of four components: true positives (TP), false positives (FP), true negatives (TN), and false negatives (FN).

Q6. Here's an example confusion matrix:

mathematica

Predicted Class	Positive	Negative	Actual Class	-----	-----
-----	Positive	TP	FN	-----	-----
				Negative	FP
					TN

From this confusion matrix, precision can be calculated as  $TP / (TP + FP)$ , recall as  $TP / (TP + FN)$ , and the F1 score as  $2 * (precision * recall) / (precision + recall)$ .

Q7. Choosing an appropriate evaluation metric for a classification problem is crucial because different metrics focus on different aspects of model performance. This can be done by considering the specific goals and requirements of the application. For example, if the cost of false positives and false negatives is asymmetric, then metrics like precision, recall, or the F1 score may be more suitable.

Q8. An example of a classification problem where precision is the most important metric is spam email detection. In this case, false positives (classifying a legitimate email as spam) are more costly than false negatives (classifying spam as a legitimate email) because they can lead to important emails being missed. Therefore, maximizing precision helps minimize the number of false positives.

Q9. An example of a classification problem where recall is the most important metric is medical diagnosis for a rare disease. In this scenario, false negatives (failing to diagnose a patient who actually has the disease) are more costly than false positives (misdiagnosing a healthy patient). Maximizing recall helps ensure that as many positive cases as possible are correctly identified, even if it means accepting some false positives.